

## 1. PERFORMANCE MEASURES

### 1.1 Floating Point Operations [Ex1]

Peak floating point performance (PP):

$$PP [\text{FLOP/s}] = f [\text{Hz} = \text{cycle/s}] \cdot c [\text{FLOP/cycle}] \cdot v [-] \cdot n [-]$$

$f$	core frequency in CPU cycles per second (given in Hz)
$c$	number of FLOP executed in each cycle
$v$	SIMD width (in number of floats)
$n$	number of cores (#cores/die * #dies/node)

### 1.2 System Memory Bandwidth [Ex1]

Theoretical memory bandwidth (PB) for DRAM:

$$PB [\text{B/s}] = f [\text{Hz} = \text{cycle/s}] \cdot c [\text{channel}] \cdot l [\text{lane/channel}] \cdot b [\text{bit/lane/cycle}] \cdot 0.125 [\text{B/bit}]$$

$f$	DRAM frequency
$c$	number of memory channels
$l$	width of the memory channel (typically 64 lanes per channel)
$b$	number of bits moved through a lane in each cycle

If there is more than one die in the node,  $PB$  has to be multiplied with the number of dies in the node.

## 2. BRUTUS

Login:

```
> ssh username@brutus.ethz.ch
```

Load/Unload a module:

```
> module load <modulename>
> module unload <modulename>
```

List all currently loaded modules:

```
> module list
```

List all available modules:

```
> module avail
```

Submit a job:

```
> bsub -n <#cores> -W <wall_clock_time> -o <output_file>
    <program_name> <program_args>
```

For example (2 hours wall clock time):

```
> bsub -n 48 -W 02:00 -o output.log ./myprogram 0 1
```

Get status of jobs:

```
> bjobs
```

## 3. MONTE CARLO ESTIMATOR

### 3.1 Formula [Ex2]

In Monte Carlo simulations we generate a set of random variables  $X_i = f(x_i)$ . From this sample we want to extract good estimates of the properties of the function  $f$ . In the lecture it was shown that the sample mean  $\bar{X}$  is a true (i.e. unbiased) estimator for the expectation value  $E[X]$  and mentioned that the variance can be estimated from the sample mean of squares.

Show that the formula given in the lecture is an unbiased estimator, i.e. prove that

$$E\left[\frac{N}{N-1}(\bar{X}^2 - \bar{X}^2)\right] = \text{Var } X$$

Proof: Let  $i = \{1, \dots, N\}$ . Insert the definition of the mean:

$$E\left[\frac{N}{N-1}(\bar{X}^2 - \bar{X}^2)\right] = \frac{1}{N-1} E\left[\sum_i X_i^2 - \frac{1}{N} \left(\sum_i X_i\right) \left(\sum_j X_j\right)\right]$$

Split the double sum:

$$= \frac{1}{N-1} E\left[\sum_i X_i^2 - \frac{1}{N} \sum_i X_i^2 - \frac{1}{N} \sum_i \sum_{j \neq i} X_i X_j\right]$$

$X_i$  and  $X_j$  are uncorrelated for  $i \neq j$ . Hence,

$$= \frac{1}{N} \sum_i E[X_i^2] - \frac{1}{N(N-1)} \sum_i \sum_{j \neq i} E[X_i] E[X_j]$$
$$E[X^2] - E[X]^2 = \text{Var } X$$

### 3.2 Example: Monte Carlo Estimation of $\pi$ [Ex2]

In this specific case every sample  $X_i$  is either 0 or 1. Therefore,  $\sum_i X_i^2 = \sum_i X_i$ .

The error of the mean is estimated as

$$\Delta = \sqrt{\frac{\text{Var } X}{N}} = \sqrt{\frac{1}{N-1}(\bar{X}^2 - \bar{X}^2)} = \sqrt{\frac{1}{N-1}(\bar{X} - \bar{X}^2)}$$

Any stochastic error decays with  $\frac{1}{\sqrt{N}}$ .

Siehe auch ML Ex2, Q3, Seite 5 ff.

Dort wird auch das Verhalten bei niedrigen Temperaturen beschrieben.

### Example Code

```
// Exercise codes for HPC course
// (c) 2012 Jan Gukelberger, ETH Zurich
#include <random>
#include <iostream>
#include <iomanip>
#include <stdexcept>

typedef double value_type;
typedef std::size_t size_type;
typedef std::mt19937 rng_type;

value_type calcp4(rng_type& rng, size_type nsamples)
{
    std::uniform_real_distribution<value_type> dist(0,1);
    size_type hits = 0;
    for( size_type i = 0; i < nsamples; ++i )
    {
        value_type x = dist(rng);
        value_type y = dist(rng);
        hits += ( (x*x + y*y) < 1. );
    }

    return hits / value_type(nsamples);
}

int main(int argc, const char** argv)
{
    if( argc != 3 )
        throw std::runtime_error(std::string("usage: ")
                                   + argv[0] + " NUM_SAMPLES SEED");

    const size_type nsamples = std::stoul(argv[1]);
    const size_type seed = std::stoul(argv[2]);
    std::cout << "threads=1, #samples=" << nsamples
               << ", seed=" << seed << ": " << std::flush;

    rng_type rng(seed);
    value_type mean = calcp4(rng, nsamples);

    // evaluate results:
    // As X_i \elem {0,1}, we don't need to accumulate the sum
    // of squares explic. but have \overline{X^2}=\overline{X}
    value_type error = std::sqrt(1./(nsamples-1.) *
                                   (mean - mean*mean));

    std::cout << std::setprecision(10) << "pi = " << 4*mean
               << " +/- " << 4*error << std::endl;
}
```

## 4. Error Estimation Code

```
#include <cmath>

class accumulator {
public:
    typedef double value_type;
    typedef std::size_t size_type;

    accumulator() : count_(0), sum_(0.), sum2_(0.) {}

    void operator<< (value_type x)
    {
        sum_ += x;
        sum2_ += x*x;
        ++count_;
    }

    void merge(accumulator const& acc)
    {
        sum_ += acc.sum_;
        sum2_ += acc.sum2_;
        count_ += acc.count_;
    }

    value_type mean() const
    {
        return sum_ / count_;
    }

    value_type error() const
    {
        return std::sqrt( (sum2_ - sum_*sum_/count_)
                          / (count_*(count_-1)) );
    }
private:
    size_type count_;
    value_type sum_, sum2_;
};
```

## 5. TIME AND MISCELLANEOUS STUFF

Method using C++11 chrono:

```
#include <chrono>
std::chrono::time_point<std::chrono::high_resolution_clock> start, end;
start = std::chrono::high_resolution_clock::now();
end = std::chrono::high_resolution_clock::now();
int elapsed_time =
std::chrono::duration_cast<std::chrono::microseconds>(end-
                                                         start).count();
```

Method using lrt:

```
#include <time.h>
timespec timeStart, timeFinish;
clock_gettime(CLOCK_MONOTONIC, &timeStart);
clock_gettime(CLOCK_MONOTONIC, &timeFinish);
double duration = timeFinish.tv_sec +
                  (double)timeFinish.tv_nsec/1e9 -
                  timeStart.tv_sec -
                  (double)timeStart.tv_nsec/1e9;
```

Get UNIX time:

```
#include <time.h>
time_t seconds;
seconds = time(NULL);
```

**Linux: Dateien durchsuchen**

```
find . -type f -name *.html -exec grep -l -i "suchstr" '{}' \;
```

## 6. ATOMIC BARRIER

```
#ifndef HPC12_ATOMICBARRIER_HPP
#define HPC12_ATOMICBARRIER_HPP

#include <atomic>
#include <cassert>
#include <limits>

class atomicbarrier
{
public:
    atomicbarrier(unsigned int count)
    : m_total(count)
    , m_count(count)
    , m_generation(0)
    {
        assert(count != 0);
    }

    void wait()
    {
        // needs to be sequentially consistent
        unsigned int gen = m_generation.load();

        // decrease the count
        // fetch_sub returns the value *before* the operation
        if (m_count.fetch_sub(1)==1) { // needs to be sequentially
            // consistent
            // if done reset to new generation of wait
            m_count.store(m_total,std::memory_order_release);
            m_generation.fetch_add(1,std::memory_order_release);
        }
        else
            while (gen ==
                m_generation.load(std::memory_order_relaxed))
                /* run in cricles, scream and shout */ ;

        unsigned int num_waiting() const
        {
            return m_count;
        }

private:
    unsigned int const m_total;
    std::atomic<unsigned int> m_count;
    std::atomic<unsigned int> m_generation;
};

#endif //HPC12_ATOMICBARRIER_HPP
```

## 7. MY ATOMIC BARRIER

```
#ifndef HPC12_ATOMICBARRIER_HPP
#define HPC12_ATOMICBARRIER_HPP

#include <atomic>
#include <cassert>

class atomicbarrier
{
public:
    atomicbarrier(unsigned int count)
    : a_total(count)
    {
        assert(count != 0);
        a_count = ATOMIC_VAR_INIT(a_total);
        a_generation = ATOMIC_VAR_INIT(0);
    }

    void wait()
    {
        unsigned int gen = a_generation;

        // decrease the count
        if (--a_count==0) {
            // if done reset to new generation of wait and wake up
            // all threads
            a_count = ATOMIC_VAR_INIT(a_total);
            a_generation++;
        }
        else
            while (gen == a_generation) {}
    }

private:
    std::atomic<unsigned int> a_count;
    unsigned int const a_total;
    std::atomic<unsigned int> a_generation;
};

#endif //HPC12_ATOMICBARRIER_HPP
```

### Barrier timing

```
#include "atomicbarrier.hpp"
#include <iostream>
#include <thread>
#include <cstdlib>
#include <vector>

int main(int argc, char** argv)
{
    // decide how many threads to use
    std::size_t nthreads =
        std::max(1u, std::thread::hardware_concurrency());
    if (argc==2)
        nthreads = std::atoi(argv[1]);
    std::cout << nthreads << " threads in total.\n";
    int repetitions=100000;
    atomicbarrier b(nthreads);

    std::vector<std::thread> threads(nthreads);

    for (unsigned i = 0; i < nthreads; ++i)
        threads[i] = std::thread( [&b,repetitions] () {
            for (int i=0;i<repetitions;++i)
                b.wait();
        });

    for (std::thread& t : threads)
        t.join();

    std::cout << "All done\n";

    return 0;
}
```

## 8. METROPOLIS PROPOSAL PROBABILITIES

Metropolis algorithm, box  $V_{\text{box}} = [0, 1]^2$  :

1. Randomly pick a particle, and propose an update  
 $\vec{x} = (x, y) \rightarrow (x + \delta x, y + \delta y)$  , where  $\delta x, \delta y$  are random (uniform) displacements in  $[-\Delta, \Delta]$ .
2. Compute the energy of the new (proposed) configuration, which is needed for the calculation of the acceptance rate  
$$P_{a,b} = \min \left[ 1, \frac{A_{b,a} w(\vec{x}_i|_b)}{A_{a,b} w(\vec{x}_i|_a)} \right], \text{ where } w(\vec{x}_i) \text{ is the Boltzmann weight.}$$
3. Accept the new configuration with probability  $P_{a,b}$ .
4. Measure the quantities you want to estimate after a fixed number of updates.

### Boltzmann weight

$$\exp \left[ -\frac{\Delta E}{kT} \right]$$

### A-priori proposal probabilities

We have to pick 1 random particle over  $N$ , and we displace it in two dimensions picking one infinitesimally small interval of size  $\delta$  over the full range  $2\Delta$ . To restore configuration  $a$  from configuration  $b$ , we need exactly the same steps:

$$A_{a,b} = A_{b,a} = \frac{1}{N} \left( \frac{\delta}{2\Delta} \right)^2$$

### Acceptance probability and detailed balance condition

The acceptance probability can be simplified to:

$$P_{a,b} = \min \left[ 1, \frac{w(\vec{x}_i|_b)}{w(\vec{x}_i|_a)} \right] = \min [1, \exp(-\beta \Delta E)] , \text{ where } \Delta E = E_b - E_a$$

The full transition rate is:

$$W_{a,b} = A_{a,b} P_{a,b}$$

We have to show that the detailed balance is satisfied, i.e.:

$$\frac{W_{a,b}}{W_{b,a}} = \frac{P_{a,b}}{P_{b,a}} = \exp(-\beta \Delta E)$$

Proof:

Suppose  $\Delta E < 0$ , then the left side of the detailed balance becomes

$$\frac{1}{\exp[-\beta(-\Delta E)]} \text{ which is equal to the right part.}$$

In the opposite case the transition rate  $W_{b,a}$  is 1, and the left part trivially becomes the right part.

## 9. BINNING ANALYSIS

→ ML Ex3 S.3 ff. Und Ex3 accumulator.hpp

## 10. MUTEX TASK QUEUE EXAMPLE

```
#include <queue>
#include <array>
#include <cmath>
#include <string>
#include <iostream>
#include <cassert>
#include <thread>
#include <mutex>
#include <condition_variable>

typedef std::size_t size_type;
typedef double value_type;
typedef std::array<value_type,2> coordinate_type;

#ifdef M_PI
#define M_PI (3.14159265358979)
#endif

// 2-dimensional trapezoidal rule for rectangular integration
// region [a,b]
template<class F>
value_type integrate(F f, const coordinate_type& a,
                    const coordinate_type& b, size_type n)
{
    ...
}

// Oscillating integrand with line singularity x^2+y^3 = 0.1
value_type integrand(const coordinate_type& x)
{
    return sin(M_PI/(x[0]*x[0]+x[1]*x[1]*x[1]-0.1));
}

struct task
{
    coordinate_type a;
    coordinate_type b;
};

struct task_queue
{
    std::mutex mutex;
    std::condition_variable available;
    std::queue<task> queue;
    size_type active = 0;

    size_type regions = 0;
    value_type result = 0.;
};
```

```
void work(task_queue& tasks, size_type n, value_type maxerror)
{
    // when there are no tasks and no active workers we are
    // finished
    std::unique_lock<std::mutex> lk(tasks.mutex);
    while( !tasks.queue.empty() || tasks.active > 0 )
    {
        // wait for available tasks
        while( tasks.queue.empty() && tasks.active > 0 )
            tasks.available.wait(lk);

        // work
        ++tasks.active;
        while( !tasks.queue.empty() )
        {
            // get task
            coordinate_type a = tasks.queue.front().a;
            coordinate_type b = tasks.queue.front().b;
            tasks.queue.pop();
            lk.unlock();

            // integrate
            value_type value = integrate(integrand,a,b,n);
            value_type error = std::abs(value -
                                      integrate(integrand,a,b,n/2));

            // subdivide into quarters by halving along each
            // dimension
            lk.lock();
            if( error > maxerror )
            {
                coordinate_type center = {{0.5*(a[0]+b[0]),
                                             0.5*(a[1]+b[1])}};
                tasks.queue.push({a,center});
                tasks.queue.push({center,b});
                tasks.queue.push({{a[0],center[1]},
                                   {center[0],b[1]}});
                tasks.queue.push({{center[0],a[1]},
                                   {b[0],center[1]}});
                tasks.available.notify_all();
            }
            else
            {
                tasks.result += value;
                ++tasks.regions;
            }
        }
        --tasks.active;
        tasks.available.notify_all();
    }
}
```

```
int main(int argc, const char** argv)
{
    // read integration parameters from command line
    if( argc != 4 )
        throw std::runtime_error(std::string("usage: ") + argv[0] +
                                   " SEGMENT_SAMPLES MAX_ERROR NUM_THREADS");
    size_type segment_samples = std::stoul(argv[1]);
    value_type max_error = std::stod(argv[2]);
    size_type nthreads = std::stoul(argv[3]);
    std::cout.precision(10);

    // full integration volume: [-1,1]^2
    task_queue tasks;
    tasks.queue.push(task({{-1,-1}}, {{1,1}} ));

    // run n worker threads
    std::vector<std::thread> workers;
    for( unsigned i = 0; i < nthreads; ++i )
        workers.push_back(std::thread(std::bind(work,
                                                  std::ref(tasks), segment_samples, max_error)));
    for( std::thread& t : workers )
        t.join();

    std::cout << "SEGMENT_SAMPLES = " << segment_samples
               << ", MAX_ERROR = " << max_error
               << ", NUM_THREADS = " << nthreads
               << ", # Regions = " << tasks.regions
               << ", Result = " << tasks.result << std::endl;
}
```

## 10. OPENMP MONTE CARLO

```
value_type mean = 0;
#pragma omp parallel num_threads(nthreads) reduction(+:mean)
{
    assert( nthreads ==
           static_cast<size_type>(omp_get_num_threads()) );
    // seed RNG with unique ID
    rng_type rng(nthreads*seed+omp_get_thread_num());
    mean = calcp4(rng, nsamples/double(nthreads)+0.5);
}
```

## 11. OPENMP N-BODY METROPOLIS

```
double calculate_energy() const
{
    double energy = 0.0;
#pragma omp parallel for schedule(guided) reduction(+:energy)
    for(size_type i=0; i < n_; ++i) {
        double energy_part = 0.0;
        for(size_type j=0; j < i; ++j)
            energy_part += potential_(config_.first[i]-
                                     config_.first[j], config_.second[i]-
                                     config_.second[j]);
        energy += energy_part;
    }
    return energy;
}

double calculate_energy_diff(size_type p,
                             std::pair<double,double> const& new_pos) const
{
    double dE = 0.;
#pragma omp parallel for reduction(+:dE)
    for(size_type i=0; i < n_; ++i) {
        if (i != p)
            dE += potential_(new_pos.first-config_.first[i],
                             new_pos.second-config_.second[i])
                - potential_(config_.first[p]
                             -config_.first[i], config_.second[p]
                             -config_.second[i]);
    }
    return dE;
}

#pragma omp parallel
#pragma omp single nowait
std::cout << std::endl << "Using " << omp_get_num_threads()
          << " threads." << std::endl;
```

## 12. OPENMP TASK QUEUE

```
value_type work(coordinate_type a, coordinate_type b,
                size_type n, value_type maxerror)
{
    value_type value = integrate(integrand,a,b,n);
    value_type error = std::abs(value -
                                integrate(integrand,a,b,n/2));
    // subdivide into quarters by halving along each dimension
    if( error > maxerror )
    {
        double r1, r2, r3, r4;
        coordinate_type center = {{0.5*(a[0]+b[0]),
                                   0.5*(a[1]+b[1])}};

#pragma omp task shared(r1, a, b, center, n, maxerror) untied
        r1 = work(a, center, n, maxerror);

#pragma omp task shared(r2, a, b, center, n, maxerror) untied
        r2 = work(center, b, n, maxerror);

#pragma omp task shared(r3, a, b, center, n, maxerror) untied
        r3 = work({{a[0],center[1]}}, {{center[0],b[1]}}, n,
                  maxerror);

        r4 = work({{center[0],a[1]}}, {{b[0],center[1]}}, n,
                  maxerror);

#pragma omp taskwait
        value = r1+r2+r3+r4;
    }

    return value;
}

int main(int argc, const char** argv)
{
    // read integration parameters from command line
    if( argc != 4 )
        throw std::runtime_error(std::string("usage: ") + argv[0] +
                                   " SEGMENT_SAMPLES MAX_ERROR NUM_THREADS");
    size_type segment_samples = std::stoul(argv[1]);
    value_type max_error      = std::stod (argv[2]);
    size_type nthreads        = std::stoul(argv[3]);
    std::cout.precision(10);

    value_type result;

    // full integration volume: [-1,1]^2
#pragma omp parallel num_threads(nthreads)
#pragma omp single nowait
        result = work({{-1,-1}}, {{1,1}}, segment_samples,
                      max_error);

    std::cout << "SEGMENT_SAMPLES = " << segment_samples
              << ", MAX_ERROR = " << max_error
              << ", NUM_THREADS = " << nthreads
              << ", Result = " << result << std::endl;
}
```

### 13. N-BODY LENNARD JONES

→ ML Ex4, 2

```
const unsigned DIMENSIONS = 2;

typedef std::size_t size_type;
typedef double scalar_type;
typedef std::array<scalar_type,DIMENSIONS> position;

struct potential
{
    potential(scalar_type rm, scalar_type epsilon,
              scalar_type rc=0)
    :   rm2_(rm*rm)
    ,   eps_(epsilon)
    ,   rc2_(1e10*rm)
    ,   shift_(0)
    {
        // default cut-off radius
        if(rc <= 0 )    rc = 2.5*rm/std::pow(2,1/6.);

        position x = {};
        position y = {{rc}};
        assert( x[0] == 0  && x[1] == 0 );
        assert( y[1] == 0 );
        shift_ = -(*this)(x,y,position());
        rc2_ = rc*rc;
        std::cout << "## Potential shift -V(rc=" << rc << ")="
                  << shift_ << std::endl;
    }

    // potential V(r^2)
    scalar_type operator()(scalar_type r2) const
    {
        if( r2 >= rc2_ )    return 0;

        scalar_type s2 = rm2_ / r2;
        scalar_type s6 = s2*s2*s2;
        return eps_*(s6*s6 - 2*s6) + shift_;
    }

    // potential V(x,y) considering periodic boundaries
    scalar_type operator()(const position& x, const position& y,
                           const position& extent) const
    {
        scalar_type r2 = 0.;
        for( size_type d = 0; d < DIMENSIONS; ++d )
        {
            scalar_type r = dist(x[d],y[d],extent[d]);
            r2 += r*r;
        }

        return (*this)(r2);
    }
};
```

```
/// compute the Lennard-Jones force particle y exerts on x
/// and add it to f
void add_force(position& f, const position& x,
               const position& y, const position& extent) const
{
    // distance vector r=x-y, considering periodic boundaries
    position r;
    scalar_type r2 = 0.;
    for( size_type d = 0; d < DIMENSIONS; ++d )
    {
        r[d] = dist(x[d],y[d],extent[d]);
        r2 += r[d]*r[d];
    }

    // potential cut-off
    if( r2 >= rc2_ )    return;

    // s = r_m/r
    // V(s) = eps * (s^12 - s^6)
    r2 = 1/r2;
    scalar_type s2 = rm2_ * r2; // (rm/r)^2
    scalar_type s6 = s2*s2*s2; // (rm/r)^6
    // common factor
    scalar_type fr = 12*eps_ * (s6*s6 - s6) * r2;
    for( size_type d = 0; d < DIMENSIONS; ++d )
        f[d] += fr * r[d];
    }

    scalar_type cutoff_radius() const {
        return std::sqrt(rc2_); }

private:
    scalar_type dist(scalar_type x, scalar_type y,
                    scalar_type extent) const
    {
        scalar_type r = x-y;
        if ( r < -extent/2 ) r += extent;
        else if( r > extent/2 ) r -= extent;
        return r;
    }

    scalar_type rm2_; // r_m^2
    scalar_type eps_; // \epsilon
    scalar_type rc2_; // cut-off radius r_c^2
    scalar_type shift_; // potential shift -V(r_c)
};
```

### 14. N-BODY: COMPARE MD AND MC

MD: Molecular Dynamics

MC: Monte Carlo

Taking long-time averages of the energies in the MD simulation we obtain the system's temperature

$$\frac{d}{2} k_B T = \bar{E}_{\text{kin}}$$

from the mean kinetic energy. With this we can start a MC simulation with the same parameters and compare the MC estimate  $\langle E_{\text{pot}} \rangle$  to the MD mean  $\bar{E}_{\text{pot}}$ .

## 15. N-BODY VERLET ALGORITHM

→ ML Ex4, 2

```
class simulation
{
public:
    simulation(const position& extent, const potential& pot,
               const std::vector<position>& x,
               const std::vector<position>& v )
    :   extent_(extent)
    ,   potential_(pot)
    ,   x_(x)
    ,   v_(v)
    ,   a_(x.size())
    {
        calculate_forces(a_,x_);
    }

    void evolve(scalar_type dt, size_type steps)
    {
        assert( steps >= 1 );
        configuration xold(x_), aold(a_);
        update_positions(x_,xold,v_,a_,dt);

        for( size_type s = 1; s < steps; ++s )
        {
            std::swap(x_,xold);
            #pragma omp parallel for
            for( size_type i = 0; i < x_.size(); ++i )
            {
                calculate_force(i,a_,xold);
                update_velocity(v_[i],aold[i],a_[i],dt);
                update_position(x_[i],xold[i],v_[i],a_[i],dt);
            }
            std::swap(a_,aold);
        }

        #pragma omp parallel for
        for( size_type i = 0; i < x_.size(); ++i )
        {
            calculate_force(i,a_,x_);
            update_velocity(v_[i],aold[i],a_[i],dt);
        }
    }

    void print_config() const
    {
        for( size_type i = 0; i < x_.size(); ++i )
            std::cout << x_[i] << v_[i] << a_[i] << std::endl;
        std::cout << std::endl;
    }
}
```

```
std::pair<scalar_type,scalar_type> measure_energies() const
{
    scalar_type epot = 0, ekin = 0;
    #pragma omp parallel for reduction(+:epot,ekin)
    for( size_type i = 0; i < x_.size(); ++i )
    {
        const position& xx = x_[i];
        const position& vv = v_[i];

        ekin += std::inner_product(vv.begin(),vv.end(),
                                   vv.begin(),scalar_type(0));

        for( size_type j = 0; j < i; ++j )
            epot += potential_(xx,x_[j],extent_);
    }

    return std::make_pair(0.5*ekin,epot);
}

private:
    typedef std::vector<position> configuration;

    void update_positions(configuration& x,
                          const configuration& xold, const configuration& v,
                          const configuration& a, scalar_type dt)
    {
        #pragma omp parallel for
        for( size_type i = 0; i < x.size(); ++i )
            update_position(x[i],xold[i],v[i],a[i],dt);
    }

    void update_velocities(configuration& v,
                           const configuration& aold,
                           const configuration& a, scalar_type dt)
    {
        #pragma omp parallel for
        for( size_type i = 0; i < v.size(); ++i )
            update_velocity(v[i],aold[i],a[i],dt);
    }

    void calculate_forces(configuration& a,
                           const configuration& x)
    {
        #pragma omp parallel for
        for( size_type i = 0; i < x.size(); ++i )
            calculate_force(i,a,x);
    }

    void update_position(position& xx, const position& xxold,
                          const position& vv, const position& aa, scalar_type dt)
    {
        for( size_type d = 0; d < DIMENSIONS; ++d )
        {
            // Verlet step
            xx[d] = xxold[d] + vv[d]*dt + 0.5*dt*dt*aa[d];

            // enforce periodic boundaries
            xx[d] = fmod(xx[d],extent_[d]);
            if( xx[d] < 0 )    xx[d] += extent_[d];
            assert( xx[d] >= 0 && xx[d] < extent_[d] );
        }
    }
}
```

```
void update_velocity(position& vv, const position& aold,
                     const position& aa, scalar_type dt)
{
    for( size_type d = 0; d < DIMENSIONS; ++d )
        vv[d] += 0.5*dt*(aold[d] + aa[d]);
}

void calculate_force(size_type i, configuration& a,
                     const configuration& x)
{
    const position& xx = x[i];
    position& aa = a[i];
    std::fill(aa.begin(),aa.end(),scalar_type(0));

    for( size_type j = 0; j < x.size(); ++j )
    {
        if( j == i ) continue;
        potential_.add_force(aa,xx,x[j],extent_);
    }
}

position extent_; // system extent along each dimension
potential potential_;

configuration x_;
configuration v_;
configuration a_;
};
```



## 16. DIFFUSION 2D

Diffusion equation:  $\frac{\partial \rho(\mathbf{r}, t)}{\partial t} = D \nabla^2 \rho(\mathbf{r}, t)$

2D:  $\frac{\partial f}{\partial \rho} = D \Delta \rho = D \nabla^2 \rho = D \frac{\partial^2 \rho}{\partial x^2} + D \frac{\partial^2 \rho}{\partial y^2}$

Discrete time:  $t = n \cdot \Delta t$

Spatial points:  $x_i = i \cdot \Delta x$  ,  $y_j = j \cdot \Delta y$

Forward-Euler:  

$$\frac{\rho_{i,j}^{(n+1)} - \rho_{i,j}^{(n)}}{\Delta t} = D \left( \frac{\rho_{i+1,j}^{(n)} - 2\rho_{i,j}^{(n)} + \rho_{i-1,j}^{(n)}}{\Delta x^2} + \frac{\rho_{i,j+1}^{(n)} - 2\rho_{i,j}^{(n)} + \rho_{i,j-1}^{(n)}}{\Delta y^2} \right)$$

To analyse the stability we perform an von Neumann stability analysis, using a plane wave ansatz:

$$\rho_{i,j}^n = \zeta^n e^{i(k_x x + k_y y)}$$

to solve the differential equation. For stability we require

$$|\zeta| \leq 1$$

otherwise the solution will grow and diverge.

From solving the discretised differential equation using our plane wave ansatz we obtain

$$\begin{aligned} \zeta &= 1 - \Delta t \left( \frac{e^{ik_x \Delta x} - 2 + e^{-ik_x \Delta x}}{\Delta x^2} + \frac{e^{ik_y \Delta y} - 2 + e^{-ik_y \Delta y}}{\Delta y^2} \right) \\ &= 1 - \left[ \frac{D \Delta t}{\Delta x^2} \right] (2 \cos(k_x \Delta x) - 2) + \left[ \frac{D \Delta t}{\Delta y^2} \right] (2 \cos(k_y \Delta y) - 2) \\ &= 1 - 4 \left[ \frac{D \Delta t}{\Delta x^2} \right] \sin^2 \left( \frac{k_x \Delta x}{2} \right) - 4 \left[ \frac{D \Delta t}{\Delta y^2} \right] \sin^2 \left( \frac{k_y \Delta y}{2} \right) \end{aligned}$$

which leads to the stability condition

$$4 \frac{D \Delta t}{\Delta x^2} + 4 \frac{D \Delta t}{\Delta y^2} \leq 2$$

The time step should therefore satisfy

$$\Delta t \leq \frac{\Delta x^2 \Delta y^2}{(\Delta x^2 + \Delta y^2) 2D}$$

For equal spacing  $\Delta x = \Delta y$  this simplifies to

$$\Delta t \leq \frac{\Delta x^2}{4D}.$$

## 17. COMPLETE MPI EXAMPLE

```
// Example solutions for HPC course
// (c) 2012 Jan Gukelberger, ETH Zurich
// (c) 2012 Michele Dolfi, ETH Zurich
```

```
#include <random>
#include <iostream>
#include <iomanip>
#include <stdexcept>
#include <cassert>
#include <mpi.h>
#include <chrono>
```

```
typedef double value_type;
typedef std::size_t size_type;
typedef std::mt19937 rng_type;
```

```
value_type calcp4(rng_type& rng, size_type nsamples)
{
    std::uniform_real_distribution<value_type> dist(0,1);
    size_type hits = 0;
    for( size_type i = 0; i < nsamples; ++i )
    {
        value_type x = dist(rng);
        value_type y = dist(rng);
        hits += ( (x*x + y*y) < 1. );
    }
}
```

```
return hits / value_type(nsamples);
}
```

```
int main(int argc, char** argv)
```

```
{
    MPI_Init(&argc, &argv);
    int np, rank;
```

```
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &np);
```

```
std::chrono::time_point<std::chrono::high_resolution_clock>
start, end;
if (rank==0)
    start = std::chrono::high_resolution_clock::now();
```

```
if( argc != 3 && rank == 0)
    throw std::runtime_error(std::string("usage: ") + argv[0] +
        " NUM_SAMPLES SEED");
```

```
const size_type nsamples = std::stoul(argv[1]);
const size_type seed     = std::stoul(argv[2]);
```

```
if (rank == 0)
    std::cout << "#procs=" << np << ", #samples=" << nsamples
        << ", seed=" << seed << ": " << std::flush;
```

```
rng_type rng(np*seed+rank); // seed RNG with unique ID
value_type partial = calcp4(rng,
    nsamples/double(np)+0.5);
```

```
value_type mean;
MPI_Reduce(&partial, &mean, 1, MPI_DOUBLE, MPI_SUM, 0,
    MPI_COMM_WORLD);
```

```
if (rank == 0) {
    // evaluate results:
    // As X_i \elem {0,1}, we don't need to accumulate the sum
    // of squares but have \overline{X^2}=\overline{X}
    mean /= np;
    value_type error = std::sqrt(1./(nsamples-1.) *
        (mean - mean*mean));

    std::cout << std::setprecision(10) << "pi = " << 4*mean
        << " +/- " << 4*error << std::endl;
```

```
end = std::chrono::high_resolution_clock::now();
int elapsed_time =
    std::chrono::duration_cast<std::chrono::microseconds>
        (end-start).count();

std::cout << "elapsed time: " << elapsed_time << "mus\n";
}
```

```
MPI_Finalize();
```

```
}
```

## 18. SSE

```
#include <x86intrin.h>
#include "aligned_allocator.hpp"
using hpc12::aligned_allocator;
// aligned to the cacheline size
typedef std::vector<scalar_type,
aligned_allocator<scalar_type,64> > positions_type;

// compute the Lennard-Jones force particle at position x0
scalar_type compute_force(positions_type const& positions,
                           scalar_type x0)
{
    scalar_type rm2 = rm * rm;
    scalar_type eps12 = 12*eps;
    size_type ndiv4 = N/4;

    __m128 mmx0 = _mm_load1_ps(&x0);
    __m128 mmrm2 = _mm_load1_ps(&rm2);
    __m128 mmeps12 = _mm_load1_ps(&eps12);
    __m128 mmforce = _mm_setzero_ps();

    for (size_type i=0; i<ndiv4; ++i) {
        __m128 mmxi = _mm_load_ps(&positions[i*4]);
        __m128 mmr = _mm_sub_ps(mmx0,mmxi);
        __m128 mmrinv = _mm_rcp_ps(mmr);
        __m128 mmrinv2 = _mm_mul_ps(mmrinv,mmrinv);
        __m128 mms2 = _mm_mul_ps(mmrinv2,mmrinv2);
        __m128 mms6 = _mm_mul_ps(mms2,_mm_mul_ps(mms2,mms2));

        __m128 mmpart = _mm_mul_ps(_mm_mul_ps(mmeps12,
        _mm_sub_ps(_mm_mul_ps(mms6,mms6),
        mms6)), mmrinv);

        mmforce = _mm_add_ps(mmforce, mmpart);
    }
    scalar_type forces[4];
    _mm_store_ps(forces,mmforce);

    for (size_type i=ndiv4*4 ; i<N ; ++i) {
        scalar_type r = x0 - positions[i];
        scalar_type rinvm = 1./r;
        scalar_type rinvm2 = rinvm * rinvm;
        scalar_type s2 = rm2 * rinvm2; // (rm/r)^2
        scalar_type s6 = s2*s2*s2; // (rm/r)^6
        forces[0] += 12*eps * (s6*s6 - s6) * rinvm;
    }

    return forces[0]+forces[1]+forces[2]+forces[3];
}
```

## 19. AVX

```
// compute the Lennard-Jones force particle at position x0
scalar_type compute_force(positions_type const& positions,
                           scalar_type x0)
{
    scalar_type rm2 = rm * rm;
    scalar_type eps12 = 12*eps;
    size_type ndiv8 = N/8;

    __m256 mmx0 = _mm256_set1_ps(x0);
    __m256 mmrm2 = _mm256_set1_ps(rm2);
    __m256 mmeps12 = _mm256_set1_ps(eps12);
    __m256 mmforce = _mm256_setzero_ps();

    for (size_type i=0; i<ndiv8; ++i) {
        __m256 mmxi = _mm256_load_ps(&positions[i*8]);
        __m256 mmr = _mm256_sub_ps(mmx0,mmxi);
        __m256 mmrinv = _mm256_rcp_ps(mmr);
        __m256 mmrinv2 = _mm256_mul_ps(mmrinv,mmrinv);
        __m256 mms2 = _mm256_mul_ps(mmrinv2,mmrinv2);
        __m256 mms6 = _mm256_mul_ps(mms2,
        _mm256_mul_ps(mms2,mms2));

        __m256 mmpart = _mm256_mul_ps(_mm256_mul_ps(mmeps12,
        _mm256_sub_ps(_mm256_mul_ps(mms6,mms6),
        mms6)), mmrinv);

        mmforce = _mm256_add_ps(mmforce, mmpart);
    }
    scalar_type forces[8];
    _mm256_store_ps(forces,mmforce);

    for (size_type i=ndiv8*8 ; i<N ; ++i) {
        scalar_type r = x0 - positions[i];
        scalar_type rinvm = 1./r;
        scalar_type rinvm2 = rinvm * rinvm;
        scalar_type s2 = rm2 * rinvm2; // (rm/r)^2
        scalar_type s6 = s2*s2*s2; // (rm/r)^6
        forces[0] += 12*eps * (s6*s6 - s6) * rinvm;
    }

    return std::accumulate(forces, forces+8, 0.);
}
```

## 20. SSE WITH CUTOFF

```
// compute the Lennard-Jones force particle at position x0
scalar_type compute_force(positions_type const& positions,
                           scalar_type x0, scalar_type rc)
{
    scalar_type rm2 = rm * rm;
    scalar_type eps12 = 12*eps;
    size_type ndiv4 = N/4;

    __m128 mmx0 = _mm_load1_ps(&x0);
    __m128 mmrc = _mm_load1_ps(&rc);
    __m128 mmrm2 = _mm_load1_ps(&rm2);
    __m128 mmeps12 = _mm_load1_ps(&eps12);
    __m128 mmforce = _mm_setzero_ps();
    // bit mask with 1 at the positions of the sign bits within
    // an __m128 vector
    __m128 signmask = _mm_castsi128_ps(
        _mm_set1_epi32(1u << 31));

    for (size_type i=0; i<ndiv4; ++i) {
        __m128 mmxi = _mm_load_ps(&positions[i*4]);
        __m128 mmr = _mm_sub_ps(mmx0,mmxi);
        // compute abs(r) by unsetting the sign bits
        __m128 mmabsr = _mm_andnot_ps(signmask, mmr);
        // rinvm = (abs(r) < rc) ? 1/r : 0
        __m128 mmrinv = _mm_and_ps(_mm_cmplt_ps(mmabsr,mmrc),
        _mm_rcp_ps(mmr) );
        __m128 mmrinv2 = _mm_mul_ps(mmrinv,mmrinv);
        __m128 mms2 = _mm_mul_ps(mmrinv2,mmrinv2);
        __m128 mms6 = _mm_mul_ps(mms2,_mm_mul_ps(mms2,mms2));

        __m128 mmpart = _mm_mul_ps(_mm_mul_ps(mmeps12,
        _mm_sub_ps(_mm_mul_ps(mms6,mms6),
        mms6)), mmrinv);

        mmforce = _mm_add_ps(mmforce, mmpart);
    }
    scalar_type forces[4];
    _mm_store_ps(forces,mmforce);

    for (size_type i=ndiv4*4 ; i<N ; ++i) {
        scalar_type r = x0 - positions[i];
        if (r >= rc) continue;
        scalar_type rinvm = 1./r;
        scalar_type rinvm2 = rinvm * rinvm;
        scalar_type s2 = rm2 * rinvm2; // (rm/r)^2
        scalar_type s6 = s2*s2*s2; // (rm/r)^6
        forces[0] += 12*eps * (s6*s6 - s6) * rinvm;
    }

    return forces[0]+forces[1]+forces[2]+forces[3];
}
```